

Applied Machine Learning

7/6/2023
Assignment 1
ELG5255[EG]



Group 4

1. Aya Metwally
2. Heba Mostafa
3. Amira Abu Issa

Problem 1:

First, we created the dataset as provided in text file uploaded in Brightspace with the assignment:

Code:

```
X_Train = np.array([[1.3, 3.3], [1.4, 2.5], [1.8, 2.8], [1.9, 3.1], [1.5, 1.5], [1.8, 2], [2.3, 1.9], [2.4, 1.4], [2.4, 2.4], [2.4, 3], [2.7, 2.7], [2.3, 3.2]])
Y_Train = np.array([0, 0, 0, 0, 1, 1, 1, 1, 2, 2, 2, 2])


X_Test = np.array([[1.7, 2.5], [1.9, 2.7], [2, 2.15], [2.4, 2], [2.2, 3.25], [2.4, 2.25]])
Y_Test = np.array([0, 0, 1, 1, 2, 2])
```

Output for shape:

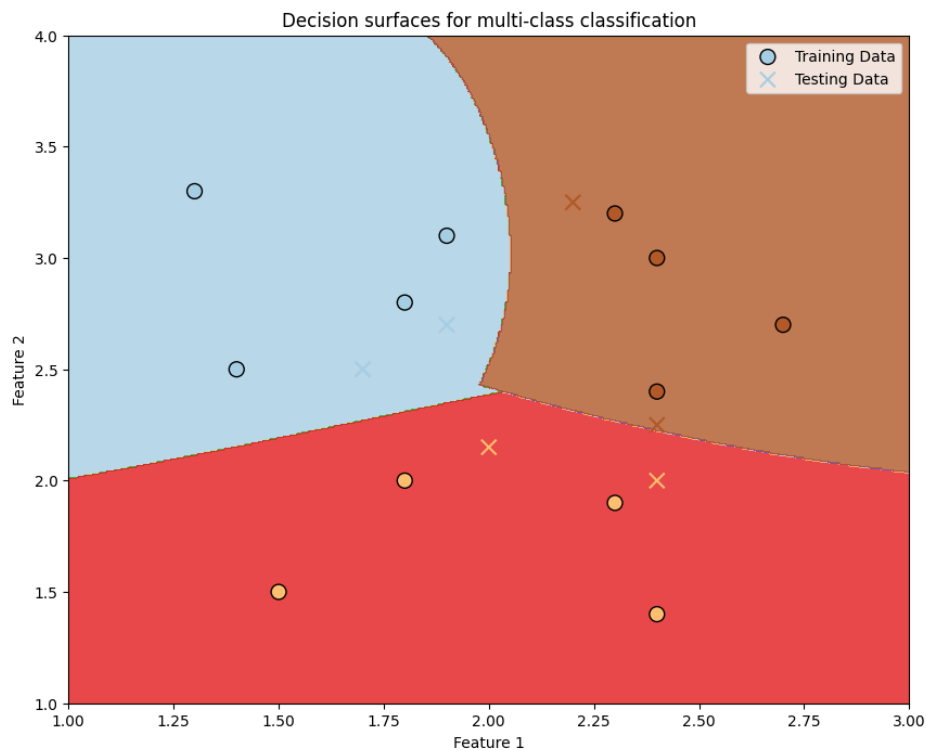

```
The shape of X-Train: (12, 2)
The shape of Y-Train: (12,)
The shape of X-Test: (6, 2)
The shape of Y-test: (6,)
```

a) Default SVM

the confusion matrix for training dataset



the confusion matrix for test dataset



b) SVM with OVR Vs. Perceptron with OVR

SVM with OVR:

for class 0:

For class : 0


Confusion matrix for training set:

```
[[8 0]
 [0 4]]
```

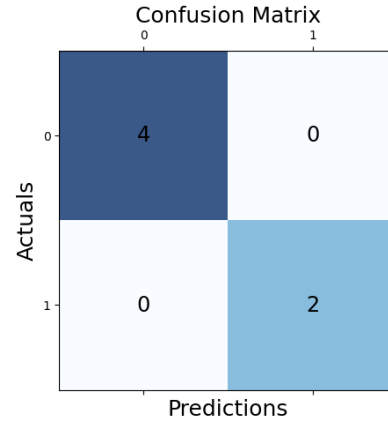
Confusion matrix for test set:

```
[[4 0]
 [0 2]]
```

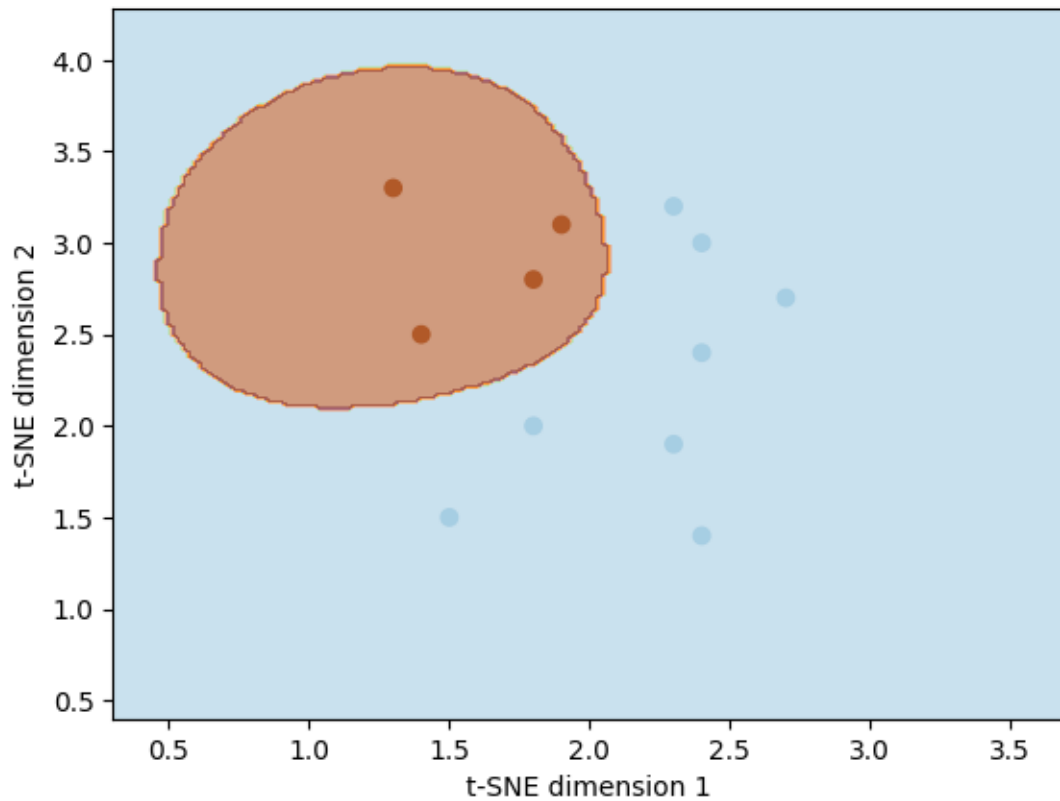
Confusion Matrix for training set:



confusion matrix for test set:




Decision boundaries



for class 1:


```
For class : 1  
Confusion matrix for training set:  
[[8 0]  
 [0 4]]
```

Confusion Matrix for training set:




```
Confusion matrix for test set:  
[[4 0]  
 [0 2]]
```

confusion matrix for test set:



Decision boundaries




for class 2:


```
For class : 2  
Confusion matrix for training set:  
[[8 0]  
 [0 4]]
```

```
Confusion matrix for test set:  
[[4 0]  
 [1 1]]
```


Confusion Matrix for training set:



confusion matrix for test set:



Decision boundaries




Perceptron with OVR:


for class 0:

```
For class: 0  
[1 1 1 1 0 0 0 0 0 0 0 0]  
[1 1 0 0 0 0]  
Confusion matrix for training dataset (one-vs-rest Perceptron):  
[[8 0]  
 [0 4]]  
Confusion matrix for test dataset (one-vs-rest Perceptron):  
[[3 1]  
 [0 2]]
```


Confusion Matrix for training set:



confusion matrix for test set:




Decision Boundaries




FOR CLASS 1:

```
For class: 1  
[0 0 0 0 1 1 1 1 0 0 0 0]  
[0 0 1 1 0 0]  
Confusion matrix for training dataset (one-vs-rest Perceptron):  
[[8 0]  
 [1 3]]  
Confusion matrix for test dataset (one-vs-rest Perceptron):  
[[4 0]  
 [1 1]]
```


Confusion Matrix for training set:



confusion matrix for test set:




Decision Boundaries





For Class 2:

```
For class: 2  
[0 0 0 0 0 0 0 1 1 1 1]  
[0 0 0 0 1 1]  
Confusion matrix for training dataset (one-vs-rest Perceptron):  
[[0 8]  
 [0 4]]  
Confusion matrix for test dataset (one-vs-rest Perceptron):  
[[0 4]  
 [0 2]]
```

Confusion Matrix for training set:



confusion matrix for test set:




SVM and Perceptron results:

SVM seeks to find the maximum margin hyperplane that separates the classes, while Perceptron seeks to find a hyperplane that minimizes the classification error. SVM can handle non-linearly separable data by using kernel functions, whereas Perceptron can only handle linearly separable data. Our dataset are non-linearly separable data, so SVM was more accurate than perceptron. SVM predicted each binary-class well, whereas perceptron didn't predict well for class c2 as shown in decision surface (the previous figure).



c) Aggregation

SVM-OVR Aggregation:

Confusion Matrix for training set:




confusion matrix for test set:




Perceptron-OVR Aggregation:


Confusion Matrix for training set:



confusion matrix for test set:



Decision Boundaries



d) Reason why SVM performance in (a) is different than aggregated performance of SVM in(c)

The aggregation in section (a) perform well with accuracy 100% than in section (c) as: In section (a), The default SVM is a multi-classification which can handle problems with more than two classes and can provide more detailed information about the data than binary classification. In section (c), the aggregation applied for binary classifiers (c0, c1, c2), where the classes are not well-separated so the accuracy of the aggregation wasn't accurate as in multi-classification as shown in the confusion matrix.

Refine the default SVM by selecting the appropriate parameter:

we select some parameters to refine the SVM model, involves kernel, C, gamma and max_iter. This can help to improve the performance of SVM and ensure that it is well-suited to the classification problem.

the code:

```
# Refine

from sklearn.model_selection import GridSearchCV

# Define the SVM classifier
clf = svm.SVC()

# Define the hyperparameters to tune
param_grid = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf', 'poly'], 'gamma': [0.1, 1, 10], 'max_iter': [100, 500, 1000]}

# Perform grid search to find the optimal hyperparameters
grid_search = GridSearchCV(clf, param_grid, cv=3)
grid_search.fit(X_Train, Y_Train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)

# Train the SVM classifier with the optimal hyperparameters
clf = svm.SVC(**grid_search.best_params_)
clf.fit(X_Train, Y_Train)


# Predict the target output for the training and test data
y_train_pred = clf.predict(X_Train)
y_test_pred = clf.predict(X_Test)

# Calculate the confusion matrices for the training and test datasets
confusion_matrix_train = confusion_matrix(Y_Train, y_train_pred)
confusion_matrix_test = confusion_matrix(Y_Test, y_test_pred)


# Print the confusion matrices
print("Confusion matrix for training dataset:\n", confusion_matrix_train)
# Display Confusion Matrix
ConfusionMatrixDisplay.from_estimator(clf, X_Train, Y_Train)
print("Confusion matrix for test dataset:\n", confusion_matrix_test)
# Display Confusion Matrix
ConfusionMatrixDisplay.from_estimator(clf, X_Test, Y_Test)
```


for class after refining:

Confusion Matrix for training set:



confusion matrix for test set:





Problem 2:

we use car-evaluation dataset, downloaded from Kaggle and implement KNN classifier and answer some questions in the assignment:

firstly, we needed to Install and import libraries:

```

▼ Install Necessary library

+ Code + Text

!pip install category_encoders

Defaulting to user installation because normal site-packages is not writeable
Collecting category_encoders
  Downloading category_encoders-2.6.1-py2.py3-none-any.whl (81 kB)
    ----- 81.9/81.9 kB 1.1 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.14.0 in c:\programdata\anaconda3\lib\site-packages (from category_encoders) (1.23.5)
Requirement already satisfied: patsy>=0.5.1 in c:\programdata\anaconda3\lib\site-packages (from category_encoders) (0.5.3)
Requirement already satisfied: statsmodels>=0.9.0 in c:\programdata\anaconda3\lib\site-packages (from category_encoders) (0.13.5)
Requirement already satisfied: pandas>=1.0.5 in c:\programdata\anaconda3\lib\site-packages (from category_encoders) (1.5.3)
Requirement already satisfied: scipy>=1.0.0 in c:\programdata\anaconda3\lib\site-packages (from category_encoders) (1.10.0)
Requirement already satisfied: scikit-learn>=0.20.0 in c:\users\dell\appdata\roaming\python\python310\site-packages (from category_encoders) (1.2.2)
Requirement already satisfied: pytz>=2020.1 in c:\programdata\anaconda3\lib\site-packages (from pandas>=1.0.5->category_encoders) (2022.7)
Requirement already satisfied: python-dateutil>=2.8.1 in c:\programdata\anaconda3\lib\site-packages (from pandas>=1.0.5->category_encoders) (2.8.2)
Requirement already satisfied: six in c:\programdata\anaconda3\lib\site-packages (from patsy>=0.5.1->category_encoders) (1.16.0)
Requirement already satisfied: joblib>=1.1.1 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (1.1.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn>=0.20.0->category_encoders) (2.2.0)
Requirement already satisfied: packaging>=21.3 in c:\programdata\anaconda3\lib\site-packages (from statsmodels>=0.9.0->category_encoders) (22.0)
Installing collected packages: category_encoders
Successfully installed category_encoders-2.6.1

▼ Import necessary libraries

[ ] import pandas as pd
    from sklearn.model_selection import train_test_split
    from sklearn.preprocessing import LabelEncoder
    from sklearn.preprocessing import OrdinalEncoder

    from sklearn.neighbors import KNeighborsClassifier
    from sklearn.model_selection import train_test_split
    from sklearn.metrics import accuracy_score
    import matplotlib.pyplot as plt
    import numpy as np
    from sklearn.model_selection import GridSearchCV
    from sklearn.metrics import classification_report
    import category_encoders as ce
  
```

a) **shuffle the dataset and split the dataset**

Shuffling and splitting the dataset is an important step, as it helps to evaluate model performance, prevent overfitting, select the best model, and avoid bias.

The code:

```
# Load the dataset
data = pd.read_csv("car_evaluation.csv")
col_names = ['Buying Price', 'Maintenance Cost', 'Doors', 'Persons', 'lug_boot', 'Safety', 'Decision']
data.columns = col_names
col_names
# Shuffle the dataset
shuffled_dataset = data.sample(frac=1, random_state=42)
print("Data After Shuffling:")
shuffled_dataset
```

The Output:

after shuffling dataset:

```
Training set = 1000
Validation set = 300
Testing set = 427

Data After Shuffling:
```

	Buying Price	Maintenance Cost	#doors	#persons	lug_boot	Safety	Decision
599	high	high	4	2	big	low	unacc
932	med	vhhigh	4	4	big	low	unacc
628	high	high	5more	2	big	high	unacc
1497	low	high	5more	4	med	med	acc
1262	med	low	4	more	med	low	unacc
...
1130	med	med	3	more	big	low	unacc
1294	med	low	5more	more	big	high	vgood
860	high	low	5more	more	big	low	unacc
1459	low	high	4	2	small	high	unacc
1126	med	med	3	more	small	high	acc

1727 rows × 7 columns

b) transform the string values into numbers

K-Nearest Neighbors (KNN) is designed to work with numerical data, not categorical data. The KNN algorithm is a type of instance-based learning, where the prediction for a new data point is based on the "k" closest training examples in feature space. The distance between the instances is computed based on the numerical values of the features. So, we transformed categorical features into numerical values before they can be used with KNN.

We use a technique called "label encoding", which assigns a unique integer value to each category. We also used Ordinal encoding, which is a technique for encoding categorical variables as numerical values based on the order or rank of the categories. We also use Ordinal encoding, a technique for encoding categorical variables as numerical values based on the order or rank of the categories.

the code:

```
#2(b)
buying = ['low', 'med', 'high', 'vhigh']
maintanance = ['low', 'med', 'high', 'vhigh']
doors = ['2', '3', '4', '5more']
persons = ['2', '4', 'more']
lug = ['small', 'med', 'big']
safety = ['low', 'med', 'high']
decision = ['unacc', 'acc', 'vgood', 'good']
all_values = [buying, maintanance, doors, persons, lug, safety]

le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_train_subset = le.transform(y_train_subset)
y_test = le.transform(y_test)
oe=OrdinalEncoder()

encoder = ce.OrdinalEncoder(cols=["buying", "maintanance", "doors", "persons", "lug_boot", "safety"])
X_train = oe.fit_transform(X_train)
X_train_subset = oe.transform(X_train_subset)
X_test = oe.transform(X_test)
```


c) different number of training samples

the code:

```
val_scores = []
test_scores = []
# Define a list of percentages of the training set to use for each classifier
train_percentages = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0]

# Iterate over each percentage and train a KNN classifier
for train_ratio in train_percentages:


    train_size = int(train_ratio * len(y_train))

    partial_X_train = X_train[:train_size]
    partial_y_train = y_train[:train_size]

    # Initialize the KNN classifier
    knn = KNeighborsClassifier(n_neighbors=2)
    knn.fit(partial_X_train, partial_y_train)
    # Train the KNN classifier on the subset of training data
    #knn.fit(X_train, y_train)
    y_valid_pred=knn.predict(X_train_subset)
    y_test_pred=knn.predict(X_test)
    # Append the accuracy scores to the respective lists
    val_accuracy = accuracy_score(y_train_subset, y_valid_pred)
    test_accuracy = accuracy_score(y_test, y_test_pred)
    val_scores.append(val_accuracy)
    test_scores.append(test_accuracy)

# Plot the accuracy scores for validation and testing sets
plt.plot(train_percentages, val_scores, label='Validation Set')
plt.plot(train_percentages, test_scores, label='Testing Set')
plt.xlabel('Percentage of Training Set')
plt.ylabel('Accuracy Score')
plt.title('Impact of Number of Training Samples on KNN Performance')
plt.legend()
plt.show()
```

the output:



The resulting plot shows how the performance of the KNN classifier varies as a function of the number of training samples used:

We can see that as we increase the number of training samples, the accuracy of the KNN classifier on the validation set and the testing set both increases. However, the rate of improvement slows down as we approach 100% of the training set. We can also see that the accuracy on the testing set is slightly lower than the accuracy on the validation set, which is expected since the model was not trained on the testing set. Overall, this plot demonstrates the importance of having a sufficient number of training samples to achieve good performance with a KNN classifier.

d) **find the best K value**

using different values of k from 1 to 10 to train KNN classifier and evaluate its performance

the code:

```
#2(d)

# Create a KNeighborsClassifier object
knn = KNeighborsClassifier()


n_neighbors = np.linspace(1, 20, 20, dtype=np.int32)
tuned_parameters = [{"n_neighbors": n_neighbors}]
knn = KNeighborsClassifier(n_neighbors=1, metric = "euclidean")
n_folds = 10
clf = GridSearchCV(knn, tuned_parameters, cv=n_folds, refit=False)
clf.fit(X_train, y_train)
scores = clf.cv_results_["mean_test_score"]
scores_std = clf.cv_results_["std_test_score"]

plt.figure().set_size_inches(8, 6)
plt.plot(n_neighbors, scores, marker="o")
std_error = scores_std / np.sqrt(n_folds)
plt.plot(n_neighbors, scores + std_error, "b--")
plt.plot(n_neighbors, scores - std_error, "b--")

# alpha=0.2 controls the translucency of the fill color
plt.fill_between(n_neighbors, scores + std_error, scores - std_error, alpha=0.2)
# Show the plot
plt.ylabel("CV score +/- std accuracy")
plt.xlabel("number of neighbors (K-Value)")
plt.axhline(np.max(scores), linestyle="--", color=".5")
plt.xticks(n_neighbors)
plt.show()

cli = KNeighborsClassifier(n_neighbors=5) ###
cli.fit(X_train, y_train)
y_pred = cli.predict(X_test)
print("Test set score: {:.2f}".format(cli.score(X_test, y_test)))
print("Final result of the model \n {}".format(classification_report(y_test, y_pred)))
```

the output:



We can see that the accuracy on the validation set is highest when $K=5$, and decreases as K increases beyond that. This demonstrates the importance of choosing an appropriate value of K for a KNN classifier to achieve good performance.

The results of

```
Test set score: 0.92
Final result of the model
```

	precision	recall	f1-score	support
0	0.84	0.80	0.82	92
1	0.91	0.59	0.71	17
2	0.94	0.99	0.96	303
3	1.00	0.60	0.75	15
accuracy			0.92	427
macro avg	0.92	0.74	0.81	427
weighted avg	0.92	0.92	0.91	427

e) **conclusions from the experiments of question (c) and (d)**

Based on the experiments conducted in questions (c) and (d), we can draw the following conclusions:

- The performance of the KNN classifier generally improves as the size of the training set increases. This is because a larger training set provides more information for the classifier to learn from and can help to reduce overfitting.
- The optimal value of K for the KNN classifier will depend on the specific dataset and problem you are working with. In our experiments in question (d), we found that the performance of the KNN classifier generally improved as K increased from 1 to around 5, and then decreased slightly as K continued to increase. However, the optimal value of K = 5

It is important to carefully choose the size of the training set and the value of K for the KNN classifier.